

# RANDOM FOREST


```
In [22]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt,seaborn as sns
```

```
In [2]: train_data=pd.read_csv(r"C:\Users\anu\Downloads\Mobile_Price_Classification_train (1).csv")
train_data
```

Out[2]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w	talk_tim
0	842	0	2.2	0	1	0	7	0.6	188	2	...	20	756	2549	9	7	1
1	1021	1	0.5	1	0	1	53	0.7	136	3	...	905	1988	2631	17	3	
2	563	1	0.5	1	2	1	41	0.9	145	5	...	1263	1716	2603	11	2	
3	615	1	2.5	0	0	0	10	0.8	131	6	...	1216	1786	2769	16	8	1
4	1821	1	1.2	0	13	1	44	0.6	141	2	...	1208	1212	1411	8	2	1
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
1995	794	1	0.5	1	0	1	2	0.8	106	6	...	1222	1890	668	13	4	1
1996	1965	1	2.6	1	0	0	39	0.2	187	4	...	915	1965	2032	11	10	1
1997	1911	0	0.9	1	1	1	36	0.7	108	8	...	868	1632	3057	9	1	
1998	1512	0	0.9	0	4	1	46	0.1	145	5	...	336	670	869	18	10	1
1999	510	1	2.0	1	5	1	45	0.9	168	6	...	483	754	3919	19	4	...

2000 rows × 21 columns



```
In [3]: test_data=pd.read_csv(r"C:\Users\anu\Downloads\Mobile_Price_Classification_test.csv")
test_data
```

Out[3]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w	talk_tim
0	1	1043	1	1.8	1	14	0	5	0.1	193	...	16	226	1412	3476	12	7	...
1	2	841	1	0.5	1	4	1	61	0.8	191	...	12	746	857	3895	6	0	
2	3	1807	1	2.8	0	1	0	27	0.9	186	...	4	1270	1366	2396	17	10	1
3	4	1546	0	0.5	1	18	1	25	0.5	96	...	20	295	1752	3893	10	0	
4	5	1434	0	1.4	0	11	1	49	0.5	108	...	18	749	810	1773	15	8	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
995	996	1700	1	1.9	0	0	1	54	0.5	170	...	17	644	913	2121	14	8	1
996	997	609	0	1.8	1	0	0	13	0.9	186	...	2	1152	1632	1933	8	1	1
997	998	1185	0	1.4	0	1	1	8	0.5	80	...	12	477	825	1223	5	0	1
998	999	1533	1	0.5	1	0	0	50	0.4	171	...	12	38	832	2509	15	11	
999	1000	1270	1	0.5	0	4	1	35	0.1	140	...	19	457	608	2828	9	2	

1000 rows × 21 columns



```
In [4]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column             Non-Null Count  Dtype  
---  -
0   battery_power      2000 non-null   int64  
1   blue                2000 non-null   int64  
2   clock_speed        2000 non-null   float64 
3   dual_sim           2000 non-null   int64  
4   fc                  2000 non-null   int64  
5   four_g             2000 non-null   int64  
6   int_memory         2000 non-null   int64  
7   m_dep              2000 non-null   float64 
8   mobile_wt          2000 non-null   int64  
9   n_cores            2000 non-null   int64  
10  pc                  2000 non-null   int64  
11  px_height           2000 non-null   int64  
12  px_width            2000 non-null   int64  
13  ram                 2000 non-null   int64  
14  sc_h                2000 non-null   int64  
15  sc_w                2000 non-null   int64  
16  talk_time           2000 non-null   int64  
17  three_g             2000 non-null   int64  
18  touch_screen        2000 non-null   int64  
19  wifi                2000 non-null   int64  
20  price_range         2000 non-null   int64  
dtypes: float64(2), int64(19)
memory usage: 328.3 KB
```

```
In [5]: test_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   id                     1000 non-null   int64  
1   battery_power          1000 non-null   int64  
2   blue                   1000 non-null   int64  
3   clock_speed            1000 non-null   float64 
4   dual_sim               1000 non-null   int64  
5   fc                     1000 non-null   int64  
6   four_g                 1000 non-null   int64  
7   int_memory             1000 non-null   int64  
8   m_dep                  1000 non-null   float64 
9   mobile_wt              1000 non-null   int64  
10  n_cores                 1000 non-null   int64  
11  pc                      1000 non-null   int64  
12  px_height              1000 non-null   int64  
13  px_width               1000 non-null   int64  
14  ram                    1000 non-null   int64  
15  sc_h                   1000 non-null   int64  
16  sc_w                   1000 non-null   int64  
17  talk_time              1000 non-null   int64  
18  three_g                1000 non-null   int64  
19  touch_screen           1000 non-null   int64  
20  wifi                   1000 non-null   int64  
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

```
In [6]: x=train_data.drop('wifi',axis=1)
        y=train_data['wifi']
```

```
In [7]: x=test_data.drop('wifi',axis=1)
        y=test_data['wifi']
```

```
In [8]: train_data['dual_sim'].value_counts()
```

```
Out[8]: dual_sim
1      1019
0       981
Name: count, dtype: int64
```

```
In [9]: test_data['dual_sim'].value_counts()
```

```
Out[9]: dual_sim
1       517
0       483
Name: count, dtype: int64
```

```
In [10]: TG={"three_g":{"Yes":1,"No":0}}
train_data=train_data.replace(TG)
print(train_data)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	842	0	2.2	0	1	0	7
1	1021	1	0.5	1	0	1	53
2	563	1	0.5	1	2	1	41
3	615	1	2.5	0	0	0	10
4	1821	1	1.2	0	13	1	44
...	...	...	...	...	...	...	...
1995	794	1	0.5	1	0	1	2
1996	1965	1	2.6	1	0	0	39
1997	1911	0	0.9	1	1	1	36
1998	1512	0	0.9	0	4	1	46
1999	510	1	2.0	1	5	1	45

	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w
0	0.6	188	2	...	20	756	2549	9	7
1	0.7	136	3	...	905	1988	2631	17	3
2	0.9	145	5	...	1263	1716	2603	11	2
3	0.8	131	6	...	1216	1786	2769	16	8
4	0.6	141	2	...	1208	1212	1411	8	2
...	...	...	...	...	...	...	...	...	...
1995	0.8	106	6	...	1222	1890	668	13	4
1996	0.2	187	4	...	915	1965	2032	11	10
1997	0.7	108	8	...	868	1632	3057	9	1
1998	0.1	145	5	...	336	670	869	18	10
1999	0.9	168	6	...	483	754	3919	19	4

	talk_time	three_g	touch_screen	wifi	price_range
0	19	0	0	1	1
1	7	1	1	0	2
2	9	1	1	0	2
3	11	1	0	0	2
4	15	1	1	0	1
...	...	...	...	...	...
1995	19	1	1	0	0
1996	16	1	1	1	2
1997	5	1	1	0	3
1998	19	1	1	1	0
1999	2	1	1	1	3

[2000 rows x 21 columns]

```
In [11]: TG={"three_g":{"Yes":1,"No":0}}
test_data=test_data.replace(TG)
print(test_data)
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	1	1043	1	1.8	1	14	0	5
1	2	841	1	0.5	1	4	1	61
2	3	1807	1	2.8	0	1	0	27
3	4	1546	0	0.5	1	18	1	25
4	5	1434	0	1.4	0	11	1	49
..	...	...	...	...	...	...	...	...
995	996	1700	1	1.9	0	0	1	54
996	997	609	0	1.8	1	0	0	13
997	998	1185	0	1.4	0	1	1	8
998	999	1533	1	0.5	1	0	0	50
999	1000	1270	1	0.5	0	4	1	35

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w
0	0.1	193	...	16	226	1412	3476	12	7
1	0.8	191	...	12	746	857	3895	6	0
2	0.9	186	...	4	1270	1366	2396	17	10
3	0.5	96	...	20	295	1752	3893	10	0
4	0.5	108	...	18	749	810	1773	15	8
..	...	...	...	...	...	...	...	...	...
995	0.5	170	...	17	644	913	2121	14	8
996	0.9	186	...	2	1152	1632	1933	8	1
997	0.5	80	...	12	477	825	1223	5	0
998	0.4	171	...	12	38	832	2509	15	11
999	0.1	140	...	19	457	608	2828	9	2

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..	...	...	...	...
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

```
In [12]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7,random_state=42)
x_train.shape,x_test.shape
```

```
Out[12]: ((700, 20), (300, 20))
```

```
In [13]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[13]: ▾ RandomForestClassifier
RandomForestClassifier()
```

```
In [14]: rf=RandomForestClassifier()
params={'max_depth':[2,3,5,10,20], 'min_samples_leaf':[5,10,20,50,100,200], 'n_estimators':[10,25,30,50,100,200]}
```

```
In [15]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')
grid_search.fit(x_train,y_train)
```

```
Out[15]: ▸ GridSearchCV
▸ estimator: RandomForestClassifier
▸ RandomForestClassifier
```

```
In [16]: grid_search.best_score_
```

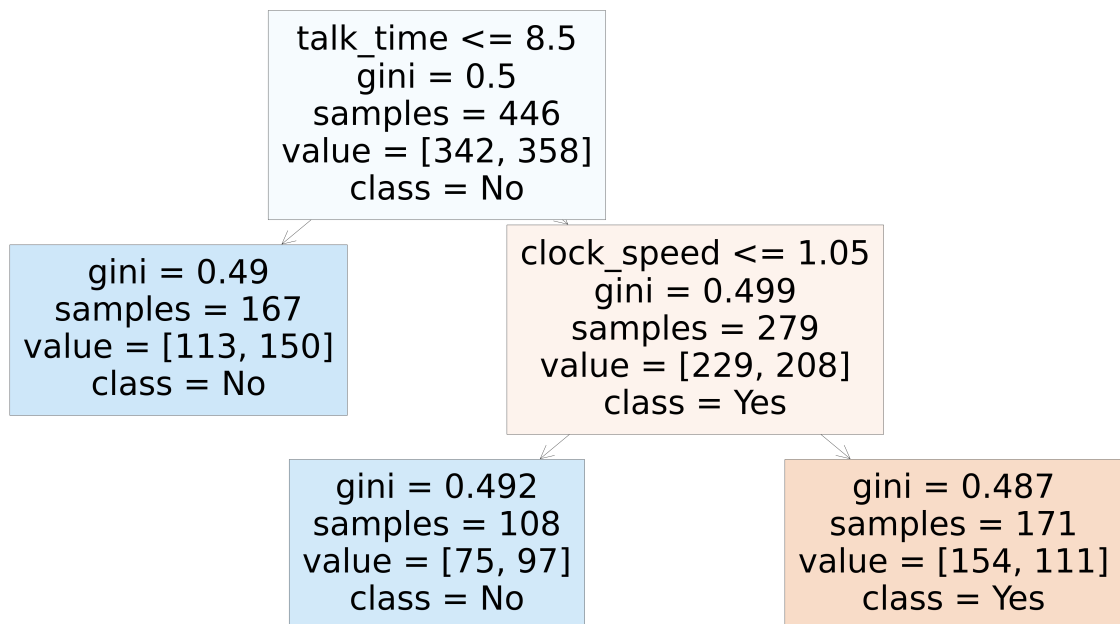
```
Out[16]: 0.5585714285714285
```

```
In [17]: rf_best=grid_search.best_estimator_
print(rf_best)

RandomForestClassifier(max_depth=20, min_samples_leaf=100, n_estimators=25)
```

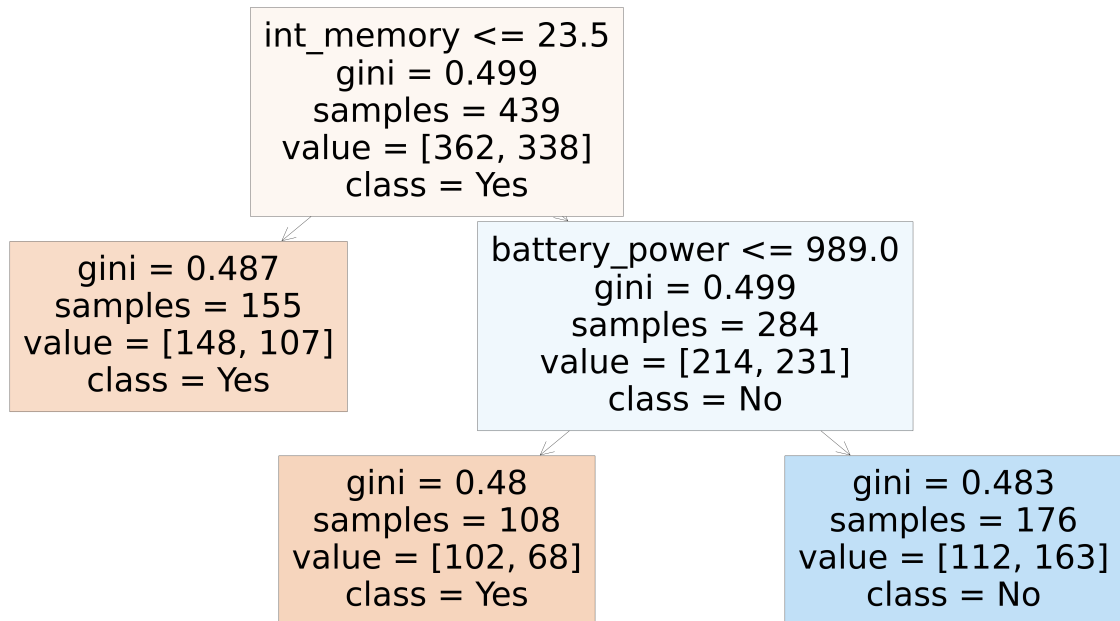
```
In [18]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=["Yes","No"],filled=True)
```

```
Out[18]: [Text(0.4, 0.8333333333333334, 'talk_time <= 8.5\ngini = 0.5\nsamples = 446\nvalue = [342, 358]\nclass = No'),
Text(0.2, 0.5, 'gini = 0.49\nsamples = 167\nvalue = [113, 150]\nclass = No'),
Text(0.6, 0.5, 'clock_speed <= 1.05\ngini = 0.499\nsamples = 279\nvalue = [229, 208]\nclass = Yes'),
Text(0.4, 0.16666666666666666, 'gini = 0.492\nsamples = 108\nvalue = [75, 97]\nclass = No'),
Text(0.8, 0.16666666666666666, 'gini = 0.487\nsamples = 171\nvalue = [154, 111]\nclass = Yes')]
```



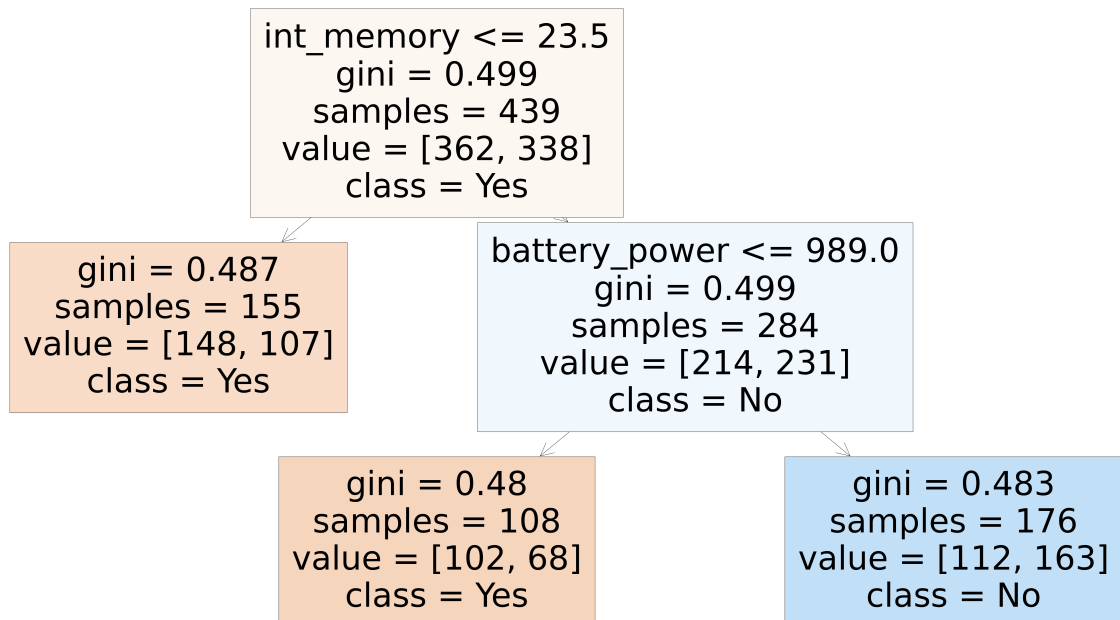
```
In [19]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=["Yes","No"],filled=True)from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=["Yes","No"],filled=True)
```

```
Out[19]: [Text(0.4, 0.8333333333333334, 'int_memory <= 23.5\ngini = 0.499\nsamples = 439\nvalue = [362, 338]\nnclass = Yes'),
Text(0.2, 0.5, 'gini = 0.487\nsamples = 155\nvalue = [148, 107]\nnclass = Yes'),
Text(0.6, 0.5, 'battery_power <= 989.0\ngini = 0.499\nsamples = 284\nvalue = [214, 231]\nnclass = No'),
Text(0.4, 0.16666666666666666, 'gini = 0.48\nsamples = 108\nvalue = [102, 68]\nnclass = Yes'),
Text(0.8, 0.16666666666666666, 'gini = 0.483\nsamples = 176\nvalue = [112, 163]\nnclass = No')]
```



```
In [20]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=["Yes","No"],filled=True)
```

```
Out[20]: [Text(0.4, 0.8333333333333334, 'int_memory <= 23.5\ngini = 0.499\nsamples = 439\nvalue = [362, 338]\nnclass = Yes'),
Text(0.2, 0.5, 'gini = 0.487\nsamples = 155\nvalue = [148, 107]\nnclass = Yes'),
Text(0.6, 0.5, 'battery_power <= 989.0\ngini = 0.499\nsamples = 284\nvalue = [214, 231]\nnclass = No'),
Text(0.4, 0.16666666666666666, 'gini = 0.48\nsamples = 108\nvalue = [102, 68]\nnclass = Yes'),
Text(0.8, 0.16666666666666666, 'gini = 0.483\nsamples = 176\nvalue = [112, 163]\nnclass = No')]
```



```
In [21]: imp_df=pd.DataFrame({"varname":x_train.columns,"Imp":rf_best.feature_importances_})
imp_df.sort_values(by="Imp",ascending=False)
```

Out[21]:

	varname	Imp
3	clock_speed	0.164792
1	battery_power	0.137961
8	m_dep	0.107444
13	px_width	0.083940
17	talk_time	0.069790
0	id	0.069107
14	ram	0.059212
7	int_memory	0.047343
9	mobile_wt	0.045225
5	fc	0.039914
11	pc	0.028644
16	sc_w	0.027381
18	three_g	0.026230
4	dual_sim	0.025145
19	touch_screen	0.021111
6	four_g	0.021096
15	sc_h	0.020025
12	px_height	0.005640
2	blue	0.000000
10	n_cores	0.000000

```
In [ ]:
```